

Taiwo Alabi

Data Scientist. I build AI systems end-to-end and publish exactly where each one fails

alabtaiwo625@gmail.com | +353 87 004 2838 | [linkedin.com/in/tami-alabi](https://www.linkedin.com/in/tami-alabi) | github.com/iamtamilore | iamtamilore.github.io | County Dublin D03 K0Y3, Ireland

Profile

I tune intelligence to **the cost of being wrong**. Every model has a dial, and the setting depends on the actual **weight of a mistake** in real life. I am a Master's student in Artificial Intelligence. I have **one product live** on the internet, as well as four **systems** I've designed and produced **end-to-end**, and I can **explain every one** of them to you during our interview.

Professional Experience

Turing, Remote (Dublin)

Jan 2026 to Mar 2026

AI Quality Analyst, Google Gemini evaluation programme

- I **measured** the quality of **Google Gemini** responses on **personalization** features, then wrote prompts to test the models, and documented the process for a **Reinforcement Learning from Human Feedback (RLHF)** project.

Huawei Technologies (via RGS), Globacom Telecommunications department, Lagos

Summer 2024 & 2025

Product Manager Intern

- Problem: On a national mobile network in one of **Africa's largest mobile markets**, **400 million daily call records represent 400 million bills** that real people pay.
- Action: I designed a **data migration architecture** that partitions customer data into **domain-specific buckets**. The system migrates **fragmented data** into **distributed read-only storage** that each team can query directly.
- Result: **Data quality improved** across the reporting platform. The **Customer Care team** can now generate **on-demand reports**, such as real-time registrations and business invoices, while query speed dropped to **under 4 seconds**.
- Together with the test team, I designed over **440 User Acceptance Test (UAT)** cases to proactively identify **faults before any system went live** or migrations occurred.

Selected Projects

Smart Medical Records Engine

FastAPI, PostgreSQL/pgvector, LangGraph, Docker

- **Problem:** A doctor has **ten minutes** and a patient history spread across **years of notes**, the one critical line, **a past reaction to a medicine**, gets missed.
- **Action:** I built a **clinical assistant** that answers **plain-English questions** from that patient's own records, with retrieval **scoped per patient** so no other patient's notes can surface, plus **physician authorization**, **PII redaction**, and an **audit log of every access**.
- **Result:** I ran the **whole stack**, the API, **PostgreSQL with 768-dimensional vector search** and a **self-hosted Llama 3.1**, across **four Docker Compose services**, so **patient data never leaves the hospital's own network**.
[Read the write-up \(PDF\)](#) | [Architecture diagram](#) | [Data model](#) | [Source code](#)

Customer Care Process Automation

Python, scikit-learn, BPMN 2.0

- **Problem:** Hundreds of customer complaints used to get **sorted by hand every morning**, which means we used to have more misrouted tickets because of **human errors**. Customers **waited for minutes** until proper routing directed their ticket to the right back-end team, followed by a few minutes of communication with that customer to resolve the issue.
- **Action:** I modeled the process in **Business Process Model and Notation** first, then split it: a **rule-based bot** enforcing **field-level conformance** against a **standard report catalogue** (the kind of signed field definitions I worked to at Huawei), and an **AI agent** for reading **customer complaint tickets** and routing them to the appropriate back-end team, reaching a **95.2% routing accuracy** on data it had never seen.
- **Result:** I built a **confidence gate** so the agent **escalates uncertain tickets to a human** instead of guessing. On **4,000 unseen tickets**, it correctly labeled every ticket it routed and appropriately flagged the uncertain edge cases for human review.
[Read the write-up \(PDF\)](#) | [BPMN process models](#) | [Confusion matrix](#) | [Source code](#)

TermsGuard AI

Sentence transformers, FAISS, Claude API, Flask

- **Problem:** A 2008 study estimated that reading all the privacy policies encountered by an average internet user would take approximately **244 hours** per year, so **91% of people** agree to **data selling and tracking** without knowing it.
- **Action:** I contributed to a **retrieval-augmented pipeline** that flags **unfair clauses in plain English**, reaching 87% F1 and cutting **model hallucination by 77%** through grounding output in retrieved reference clauses.
- **Results:** I cut **end-to-end latency from 28 seconds to under 10**, and added a **Claude verification stage** that removes false positives, and built a cosine-only fallback when no API key is set.

[Live demo](#) | [Read the write-up \(PDF\)](#) | [Source code](#)

Automated QA Gate for AI Generated Imagery

TensorFlow/Keras, ResNet50, scikit-learn

- **Problem:** **AI-generated fashion campaign images** silently drift over multiple prompt generations; the character's **identity starts to shift** or specific details I want to retain become inconsistent.
- **Action:** I built a **two-stage gate**: a **Siamese network** on a **frozen ResNet50** to flag identity drift and an **autoencoder** for style drift, reaching **precision 1.00 with zero false positives**, against pixel difference and kNN baselines at **F1 0.59 and 0.26**. Tuning the decision threshold from the sigmoid default of 0.5 to **0.4821**, on validation data only, recovered all eight missed pairs without adding a single false positive and lifted **F1 from 0.90 to 1.00**.

[Read the write-up \(PDF\)](#) | [ROC curve and confusion matrix](#) | [Architecture diagram](#) | [Source code](#)

Education

National College of Ireland, Dublin

Jan 2026 to Jan 2027

MSc Artificial Intelligence (part time)

Modules: Machine Learning, **Intelligent Agents and Process Automation**, **AI-Driven Decision Making**, **Engineering and Evaluating AI**, **Natural Language Processing**, **Data Governance and Ethics**, Database Programming, Foundations of AI.

Afe Babalola University, Lagos

Sep 2021 to Feb 2025

BSc Computer Science, Second Class Honours, Upper Division (2.1)

Technical Skills

Languages: Python, SQL

Machine Learning: Object-oriented design approaches, Agile methodologies, **scikit-learn**, **pandas**, **Matplotlib**, **transfer learning**, **reinforcement learning**, static ML, adaptive ML, Q-Learning, K-NN, Bayesian networks, **neural networks**, **Transformers (time-series)**.

AI & Agents: **RAG pipelines**, **Claude**, **MCP servers**, **LangChain**, **LangGraph**, **vector search** (pgvector, FAISS), LLM evaluation, intelligent/multi-modal agents, [genetic algorithms](#), TF-IDF vectorizer, **zero-shot learning/few-shot learning**.

Data & Databases: **PostgreSQL**, **Hadoop HDFS**, HBase, GaussDB, **data modeling**

Containers & Deployment: Docker, Docker Compose, Kubernetes, **FastAPI**, Flask, autoscaling, **Google Cloud**, Google Kubernetes Engine

Data Governance: **GDPR** data cataloging, **data lineage**, data quality validation, **UAT** design

Tools & Version Control: **Git**, GitHub, Claude Code. My server is **dual-boot (Ubuntu/Windows)**. Native **Linux** user proficient with **Microsoft 365**.